

Abhinav G Krishnan

Bangalore, India | +91 9447715481 | abhinavgkrishnan@gmail.com | [LinkedIn](#) | [Portfolio](#) | [Github](#)

Profile

AI engineer focused on turning natural-language intent into reliable, executable systems. Agent orchestration, RAG, and the eval discipline to make generated output trustworthy. Backend-strong, ships end-to-end from retrieval to production; comfortable across AI and web3.

Core Skills

- **Languages:** Python, TypeScript, JavaScript, SQL
- **AI & ML:** RAG, Vector Search, Agent Orchestration, MCP, Prompt Engineering, LLM Integration, Evals
- **Backend & Data:** FastAPI, REST APIs, Webhooks, Kafka, DuckDB, Qdrant
- **Frontend & Infra:** Next.js, Docker, GCP, Supabase, Vercel, Git

Work Experience

AI Tech Cloud Network (formerly Solidus AI Tech: AI infrastructure & automation)

AI & Fullstack Engineer

(01.02.2025 – 31.07.2026)

- Architected the backend of Forge, a visual AI workflow platform: composable engine of 18+ block types across 90+ tool integrations, MCP servers, and 6+ model providers, deployable as chat, REST API, webhook, or scheduled job.
- Built a RAG meta-agent that compiles plain-language prompts into runnable multi-step workflows, using live block metadata and a schema-validation retry loop to keep generated output valid.
- Shipped a natural-language sports-betting agent; eliminated context-overflow failures by routing oversized API responses through a DuckDB SQL layer and querying only the rows needed.
- Owned backend delivery end to end: deployment, production monitoring, Web3 wallet-connect authentication, and real-time multi-user collaboration.

Deloitte USI

Analyst

(16.01.2023 – 31.01.2025)

- Owned the NetSuite side of a Salesforce-Boomi-NetSuite integration: reconciled field mappings, proration, and numeric accuracy across all three layers, and partnered directly with client stakeholders to onboard them and triage change requests.
- Delivered a custom billing solution covering a gap in native ERP functionality.

Projects

- CompanionCube (Spoiler-aware RAG game guide): Progress-gated retrieval over 5000+ chunks (dense + BM25, RRF fusion, cross-encoder rerank) with spoiler control as vector-store payload filters; 100% recall@6, 0% leak. Extended with an LLM-extracted knowledge graph and an LLM-judge eval. Deployed on GCP Cloud Run.
- OhSnap (High-throughput Farcaster indexer): Sharded Kafka pipeline ingesting network activity in real time, with developer APIs and a live dashboard for low-latency streaming.

Education

Govt. Model Engineering College, Kochi

(01.06.2018 – 31.05.2022)

B. Tech, Electronics & Communication Engineering (First Class)